

Migrating to Vitess at (Slack) Scale

Michael Demmer

Percona Live - April 2018





slackbot just now



Welcome!

This is a (brief) story of how
Slack's databases work
today, why we're migrating
to Vitess, and some lessons
we've learned along the way.



Michael Demmer

*Senior Staff Engineer
Slack Infrastructure*

- ~1.5 years at Slack, former startup junkie
- PhD in CS from UC Berkeley
- Long time interest in distributed systems
- (Fairly) new to databases





Our Mission:
To make people's working
lives simpler, more
pleasant,
and more productive.





- 9+ million weekly active users
- 4+ million simultaneously connected
- Average 10+ hours/ weekday connected
- \$200M+ in annual recurring revenue
- 1000+ employees across 7 offices
- Customers include: Autodesk, Capital One, Dow Jones, EA, eBay, IBM, TicketMaster, Comcast



How Slack (Mostly) Works

Focusing on the MySQL parts



The Components



Linux



Apache



MySQL



HHVM



Real Time Messaging



Caching

The Components



Linux




Apache



MySQL



HHVM



Real Time Messaging



Caching


“Legacy” MySQL Numbers

Primary storage system for
the Slack service
(File uploads in AWS S3)

~1400 database hosts

~100,000-400,000 QPS
with very high bursts


~24 billion queries / day



MySQL Details

- MySQL 5.6 (Percona Distribution)
- Run on AWS EC2 instances, no containers
- SSD-based instance storage (no EBS)
- Single region, multiple Availability Zones
- Webapp has many short-lived connections directly to mysql



Master / Master



- Each is a *writable* master AND a replication slave of the other
- Fully async, statement-based replication, without GTIDs
- App *prefers* one "side" using $team_id \% 2$, switches on failure
- Mitigate conflicts by using upsert, globally unique IDs, etc
- Yes, this is a bit odd... BUT it yields Availability >> Consistency

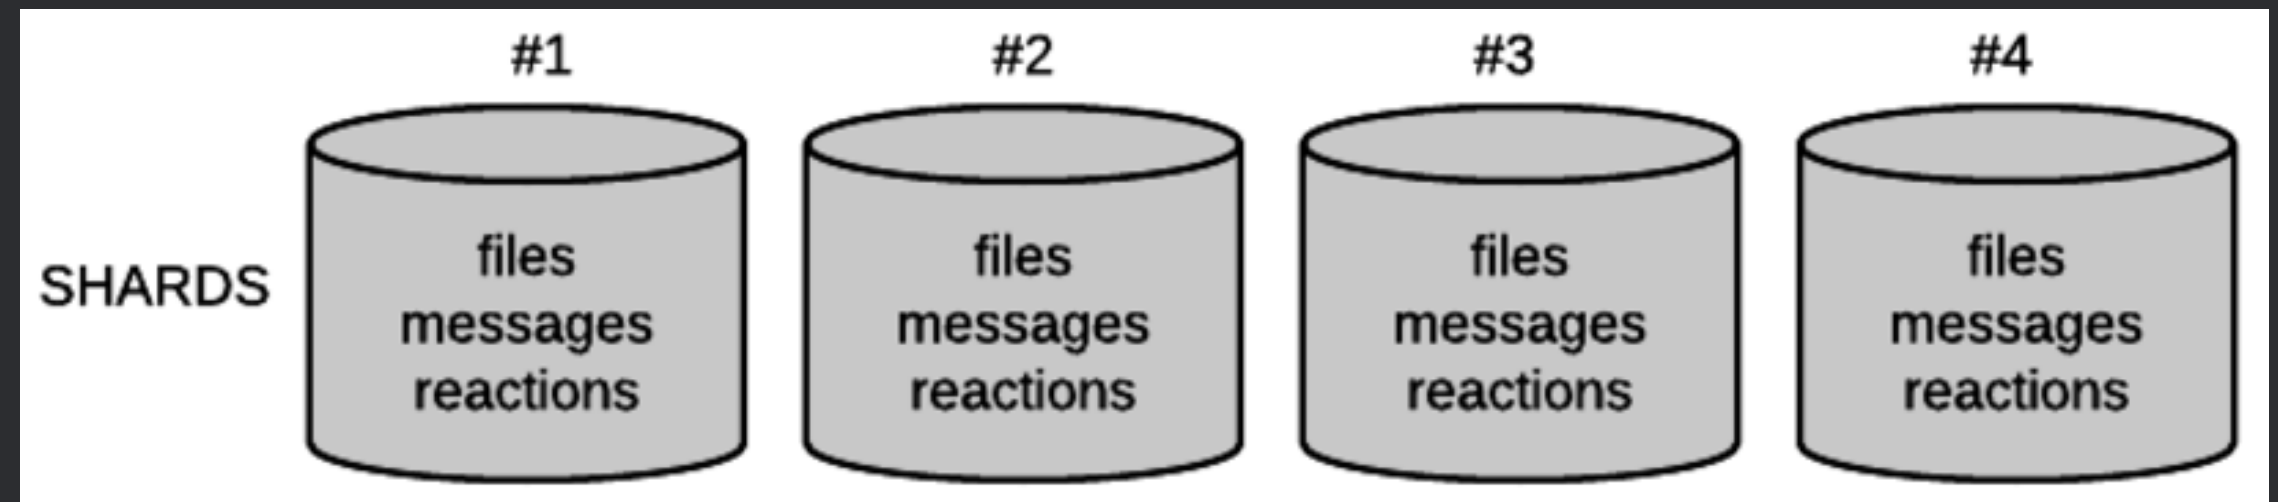
Sharding

Workspace (aka "team")
assigned to a shard at
signup



App finds team:shard
mapping in mains db


Globally Unique IDs via a
dedicated service



Added Complexity

Enterprise Grid:
Federate multiple
workspaces into an
org using $N + 1$ shards

Shared Channels:
Accessing across
workspace shards



The Good Today

- ✓ Highly available for transient or permanent host failures
- ✓ Highly reliable with low rate of conflicts in practice
- ✓ Writes are as fast as a single node can accept
- ✓ Horizontally scale by splitting "hot" shards
- ✓ Can pin large teams to dedicated hosts
- ✓ Simple, well understood, easy to administer and debug

Challenges



Hot Spots

Large customers or unexpected usage concentrates load on a single shard

Can't scale up past the capabilities of a single database host

The image displays three overlapping screenshots of GitHub commit history pages, illustrating database incidents. The top screenshot, titled "2017-08-25 shard218 thread exhaustion for one hour", shows a commit by a user who deleted a channel with many files and shares, causing a thread exhaustion on both sides of the shard for 65 minutes. The middle screenshot, titled "2017-09-07 Shards 150, 284, and 644 overwhelmed with selects", shows a commit by a user who optimized the extra ai. The bottom screenshot, titled "2017-06-28 Channel Highlights Introduced New Load on the [redacted] Shard", shows a commit by a user who introduced new load on the shard. Each screenshot includes a timeline of commits and a "Raw" button.

2017-08-25 shard218 thread exhaustion for one hour

A user deleted a channel with many files + shares. The process maxed out connections on both sides of the shard for 65 mins.

2017-09-07 Shards 150, 284, and 644 overwhelmed with selects

Timeline

- 2017-0907 12:40: @n optimize the extra ai

2017-06-28 Channel Highlights Introduced New Load on the [redacted] Shard

2017-07-13 [redacted] Flannel and Shard are Hot

Timeline

Application Complexity

Need the right context
to route a query

Scatter query to many
shards when the
“owner” team is not
known.

```
if (teams_is_on_enterprise($team)){
    $enterprise = teams_get_enterprise($team);

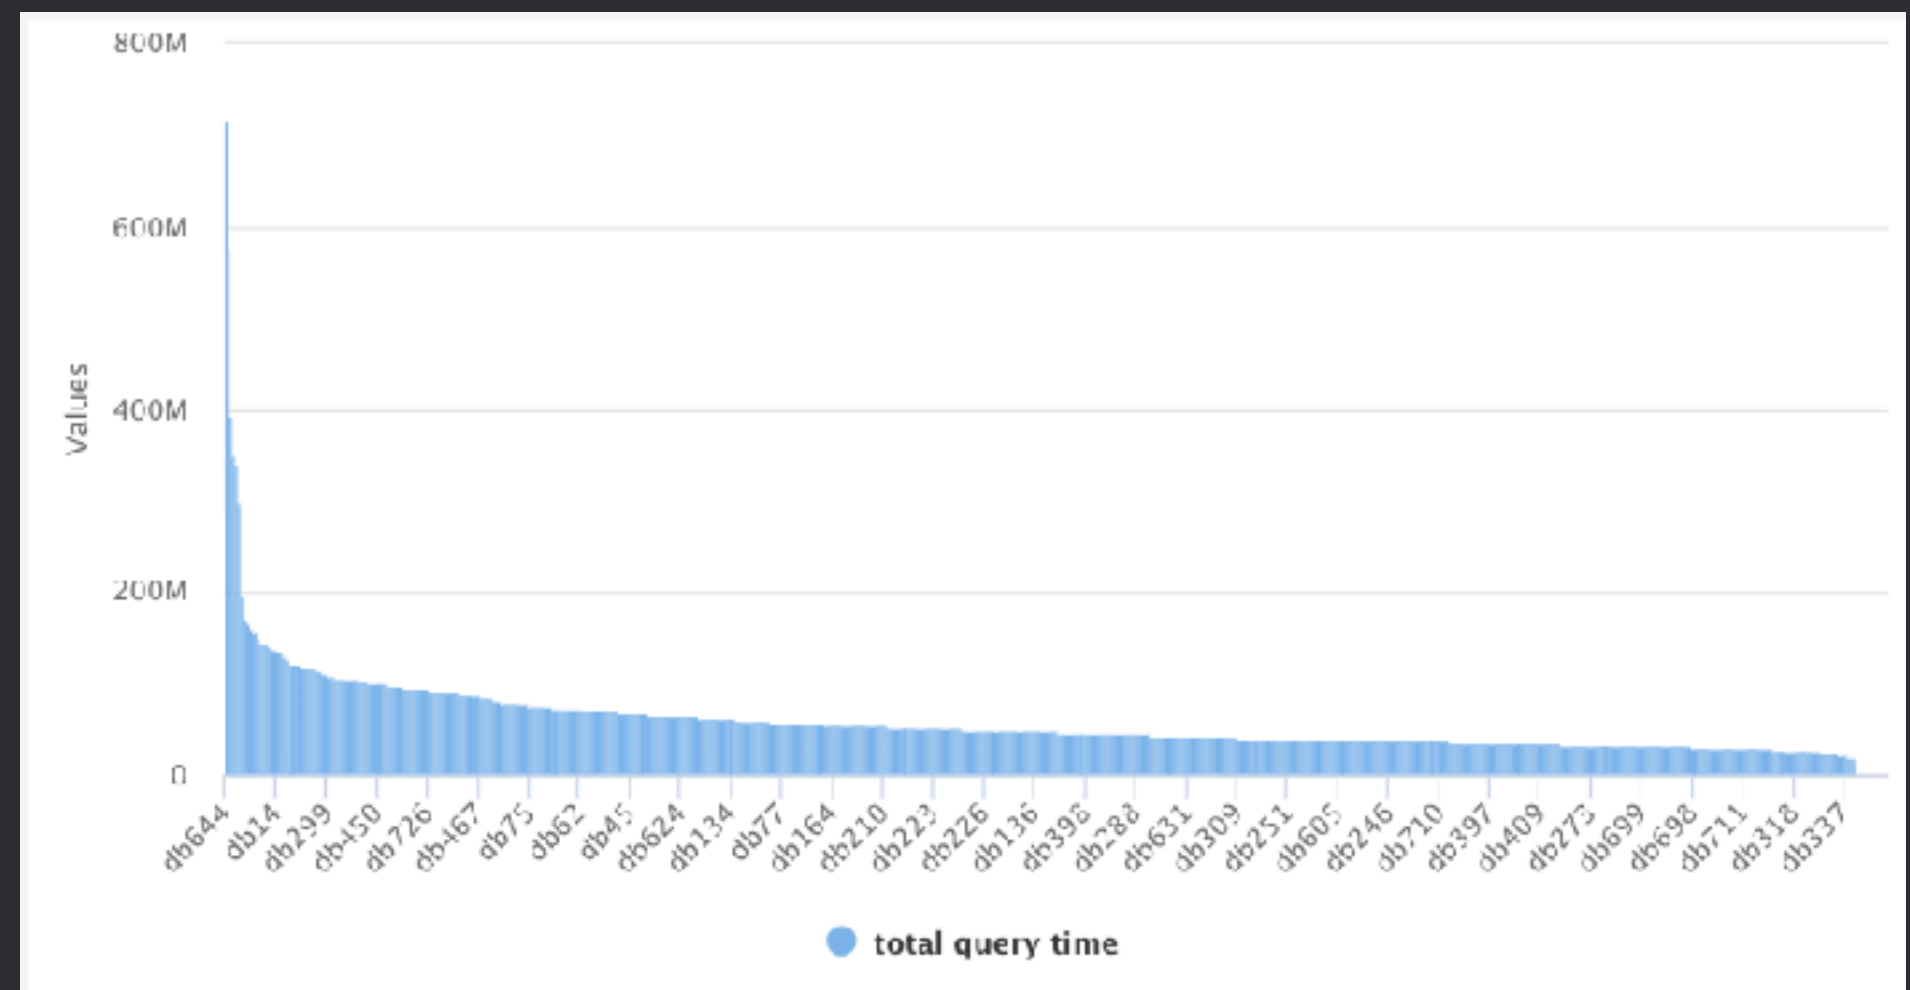
    $ret = db_fetch_team($enterprise, "SELECT * FROM teams_c
        'team_id'          => $enterprise['id'],
    ));
    foreach ($ret['rows'] as $row){
        if (!$previous_names[$row['channel_id']]) $previ
        $previous_names[$row['channel_id']][] = $row;
    }

    foreach ($previous_names as $channel_id => $names){
        usort($names, function($a, $b){ return $b['date_
        $previous_names[$channel_id] = $names;
    }
}
```

Inefficient Usage

Average load (~200 qps) much lower than capacity to handle spikes

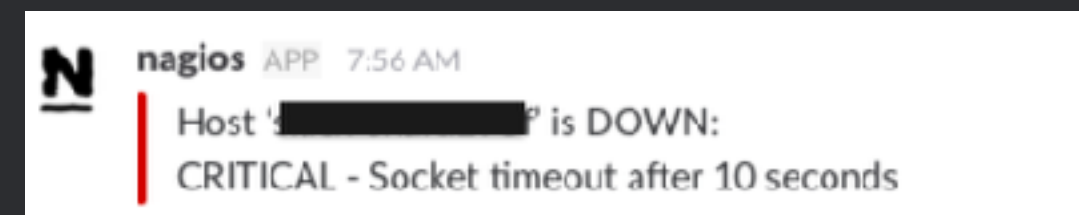
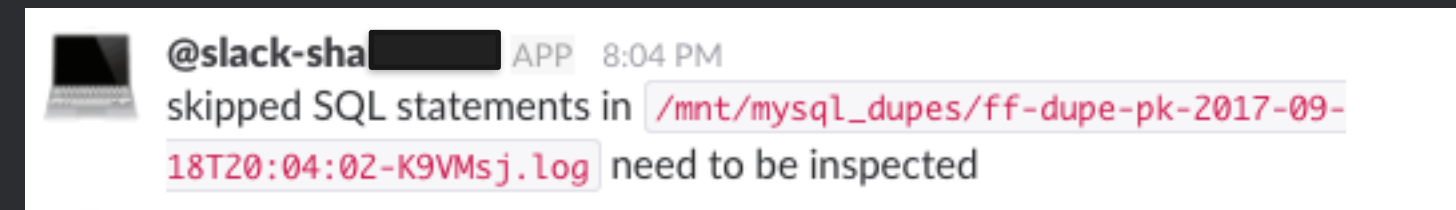
Very uneven distribution of queries across hosts



Operator Interventions

Operators need to manually repair conflicts and replace failed hosts.

Busy shards are split using manual processes and custom scripts



So What To Do?



Next Gen Database Goals

- ✨ Shard by Anything! (Channel, File, User, etc)
- 💻 Maintain Existing Development Model
- 🕒 Highly Available (but a bit more consistent)
- 📈 Efficient System Utilization
- 👌 Operable In Slack's Environment

Possible Approaches


Shard by X in PHP



- + no new components
- + easiest migration

- lots of development and operations effort

NoSQL



- + flexible sharding
- + proven at scale

- major change to app
- new operations burden

NewSQL



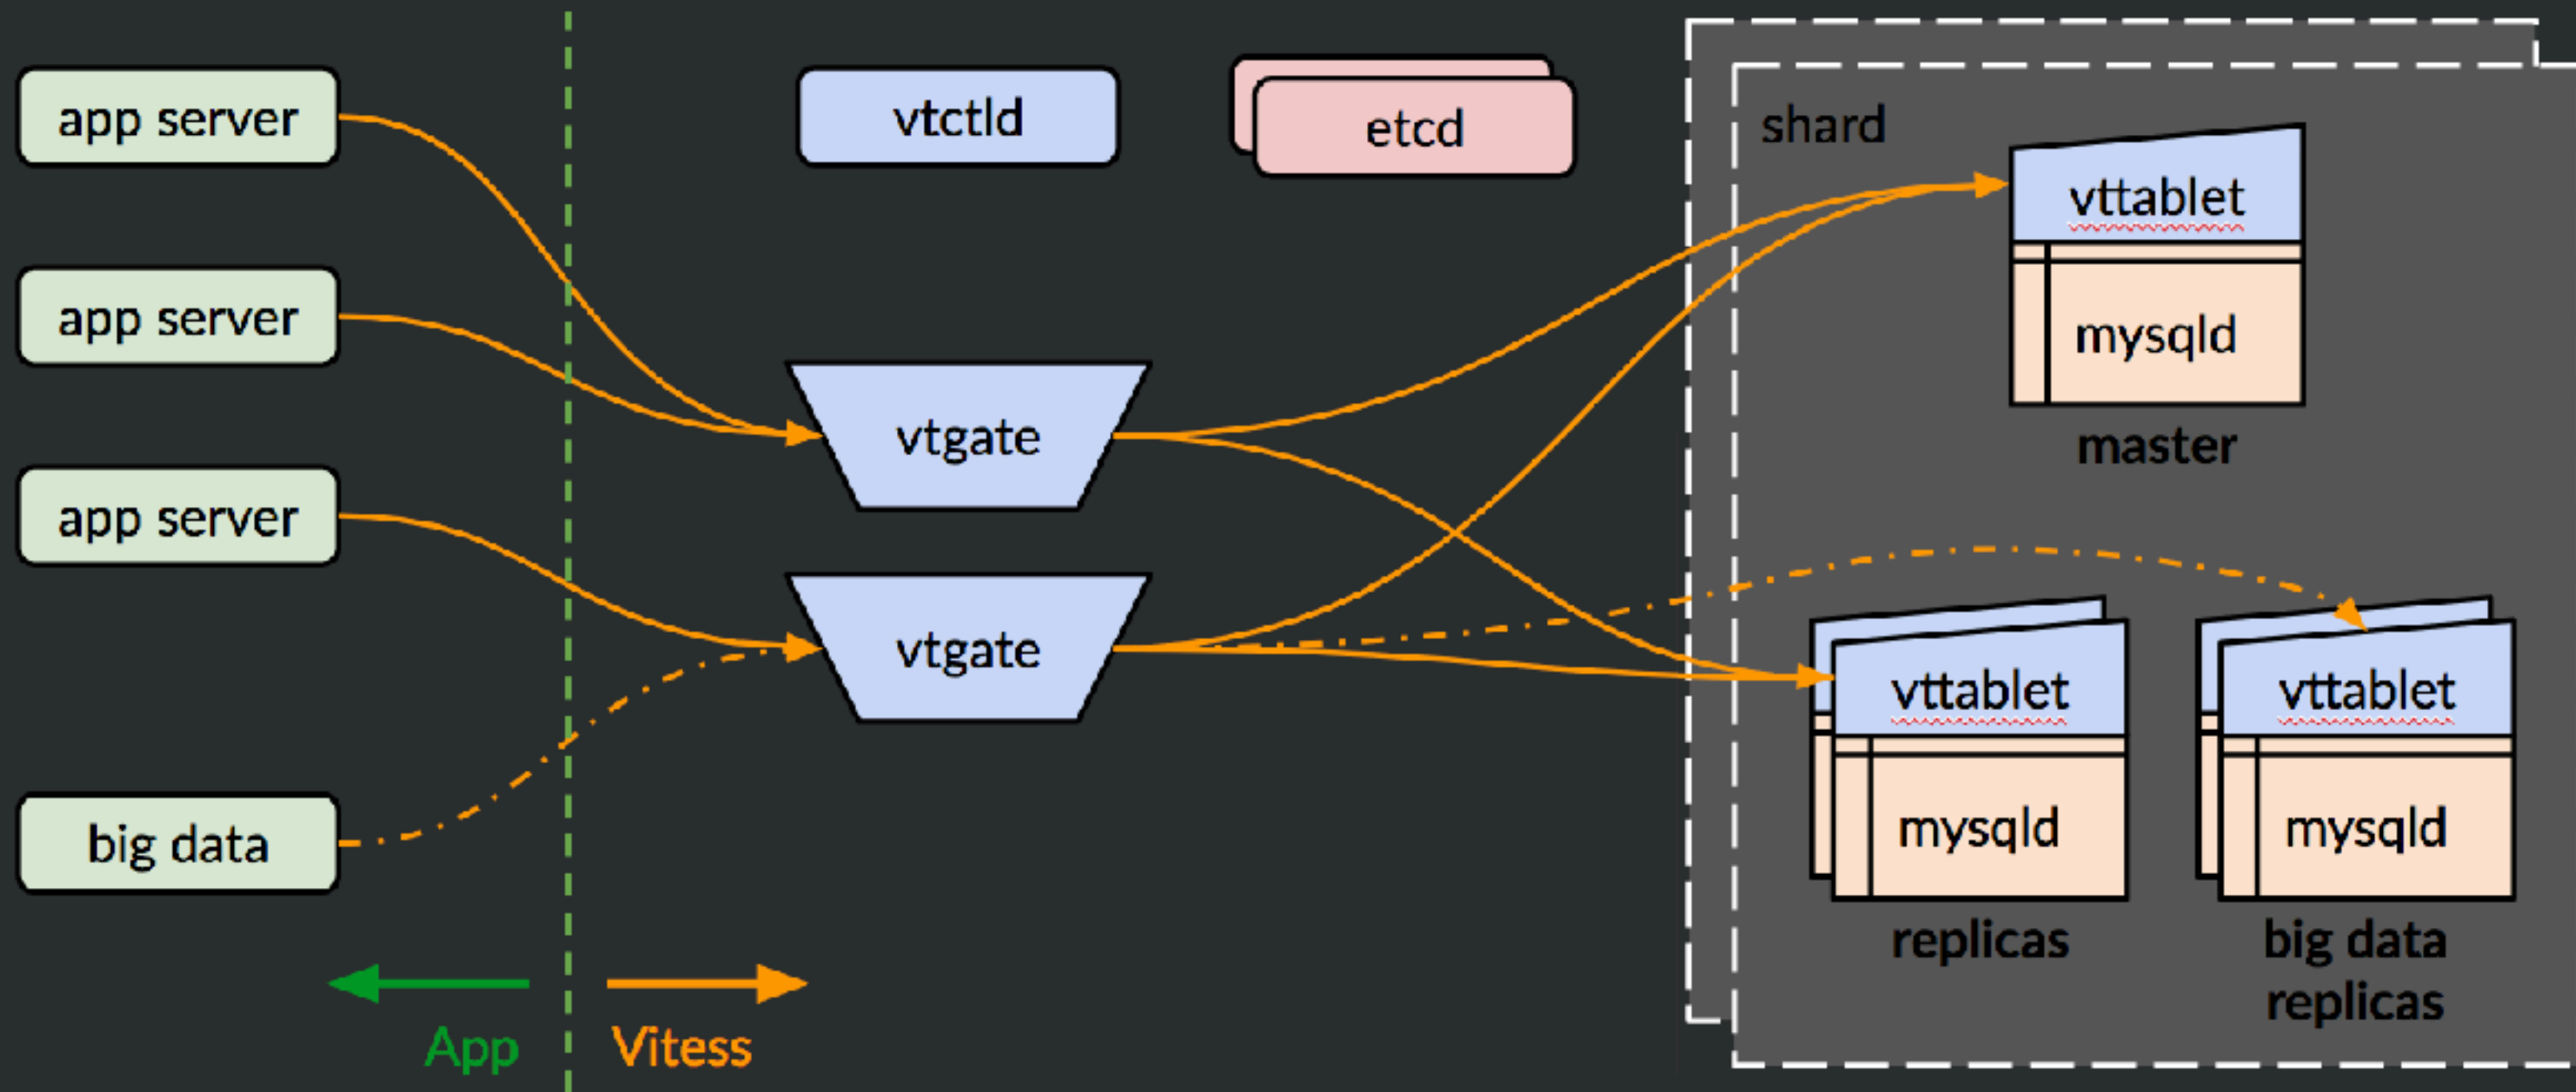
- + flexible sharding
- + scale-out storage
- + SQL compatibility!

- least well known

Why Vitess?

- Scaling and sharding flexibility without changing SQL (much)
- MySQL core maintains operator and developer know-how
- Proven at scale at YouTube and more recently others
- Active developer community and approachable code base

Vitess In One Slide



Shard by Anything

- Applications issue queries as if there was one giant database, Vtgate routes to the right shard(s)
- "Vindex" configures most natural sharding key for each table
- Aggregations / joins pushed down to MySQL when possible
- Secondary lookup indexes (unique and non-unique)
- Still supports inefficient (but rare) patterns: Scatter / gather, cross-shard aggregations / joins

Easy Development Model

- Vitess supports the mysql server protocol end to end
- App connects to any Vtgate host to access *all* tables
- Most SQL queries are supported (with some caveats)
- Additional features: connection pooling, hot row protection, introspection, metrics

Highly Available (and more consistent)

- Vitess topology manager handles master / replica config
- Actual replication still performed by MySQL
- Changed to row-based, semi-sync replication using GTIDs
- Deployed Orchestrator to manage failover in seconds

Efficient System Usage

- Vitess components are performant and well tuned from production experience at YouTube
- Can split load vertically among different pools of shards
- Even distribution of fine grained shard keys spreads load to run hosts with higher average utilization

Operable in Slack's Environment

- MySQL is production hardened and well understood
- Leverage team know-how and tooling
- Replication still uses built-in mysql support
- New tools for topology management, shard splitting / merging
- Amenable to run in AWS without containers

Vitess Adoption: Approach and Experiences



Migration Approaches

Migrate individual tables / features one by one

Run Vitess in front of existing DBs

Migration Approaches

Migrate individual tables / features one by one 

- Only approach that enables resharding (for now)
- Methodical approach to reduce risk

Run Vitess in front of existing DBs 

- Could make it work with custom sharding scheme in Vitess
- But we run master/master
- And doesn't help to avoid hot spots!

Migration Plan

- For each table to migrate:
 1. Analyze queries for common patterns
 2. Pick a keyspace (i.e. set of shards) and sharding key
 3. Double-write from the app and backfill the data
 4. Switch the app to use vitess
- But we also need to find and migrate all joined tables
 - ... and queries that aren't supported or efficient any more
 - ... and whether the old data model even makes sense!!

Offline analysis (vtexplain)

- Analysis tool to show what actually runs on each shard
- Query support is not yet (likely never be) 100% MySQL
- Choice of sharding key is crucial for efficiency

```
# vtexplain -shards 64 -schema-file test-schema.sql -vschema-file test-vschema.json -  
sql "insert into user (id, name) values (123, 'Jane Doe')"
```

```
insert into user (id, name) values (123, 'Jane Doe')
```

```
1 ks_sharded/f0-f4: begin  
1 ks_sharded/f0-f4: insert into name_user_map(name, user_id) values ('Jane Doe', 123)  
2 ks_sharded/10-14: begin  
2 ks_sharded/10-14: insert into user(id, name) values (123, 'Jane Doe')  
3 ks_sharded/f0-f4: commit  
4 ks_sharded/10-14: commit
```

Migration Stages



PASSTHROUGH: Convert call sites



BACKFILL: Double-write & bulk copy, read legacy



DARK: Double-read/write, app sees legacy results




LIGHT: Double-read/write, app sees Vitess results



SUNSET: Read/write only from Vitess

Current Status

- 🎉 Running in production for 10 months
- Serving ~10% of all queries, part of the critical path for Slack
- All new features use Vitess
- Migrating other core tables this year



Current Status: Details



- ~30,000 QPS at peak times, occasional spikes above 50,000
- 8 keyspaces, 3 replicas per shard, 316 tablets, 32 vtgates
- Query mix is ~80% read, 20% write
- Currently ~75% queries go to masters

Performance

Millisecond latencies for
connect/read/write

Slower due to extra network
hops, semi-sync waits, and
Vitess overhead

So far as expected — slightly
slower but steadier




Performance Improvements

Vitess modifications:

- Avoid round trips for autocommit transactions
- Scatter DML queries
- Query pool timeouts

Dramatically improved both average and tail latencies



Vitess Deployment: Multi AZ

us-east-1a



us-east-1b

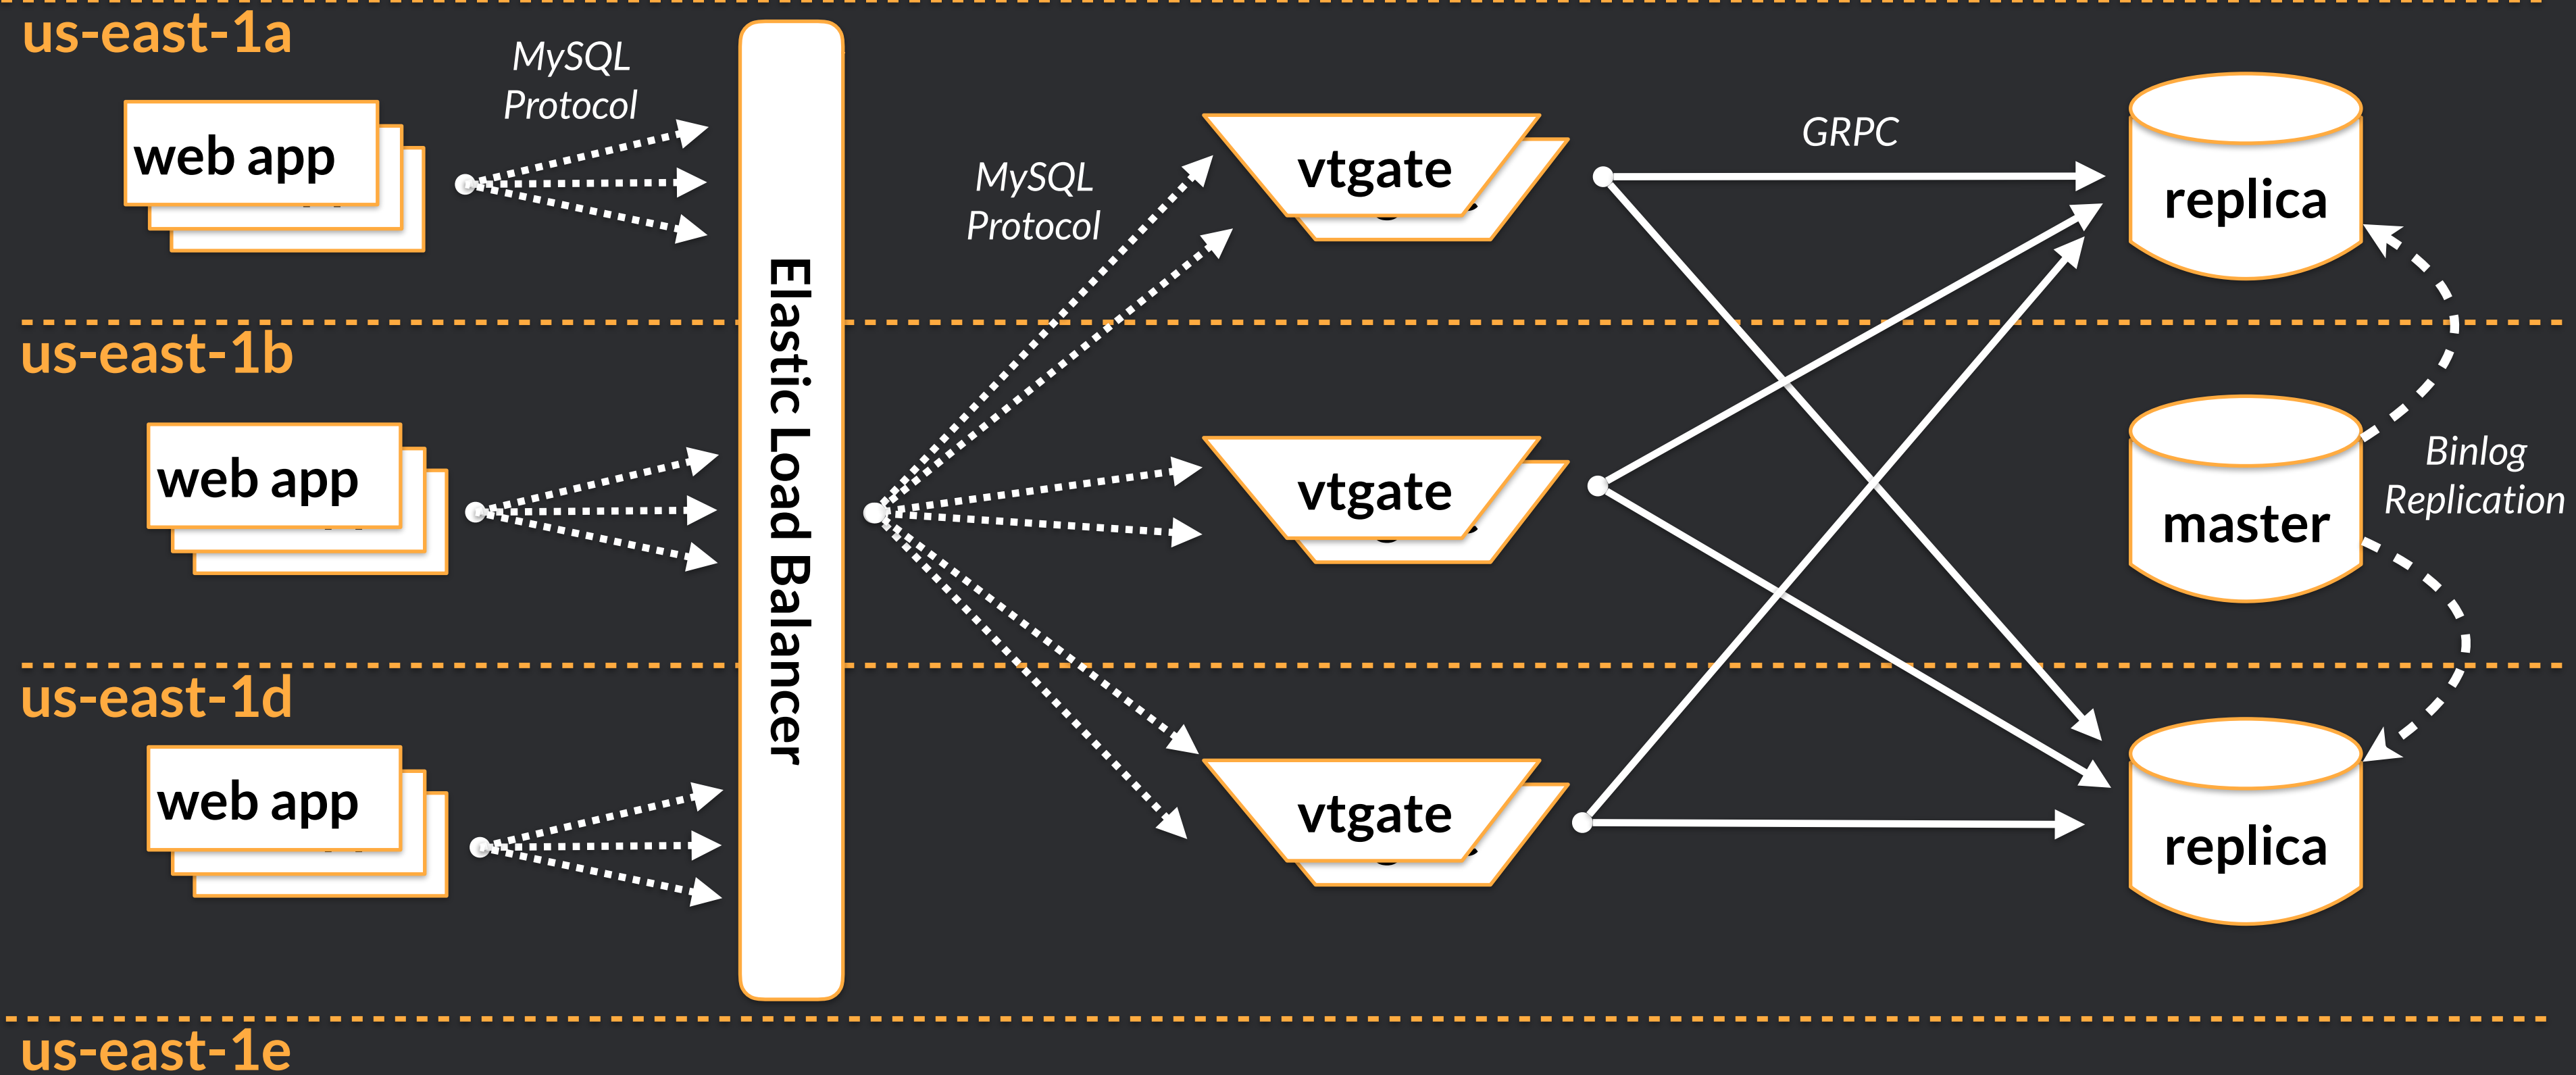


us-east-1d



us-east-1e

Initial Deployment



Client Side Load Balancing

us-east-1a



MySQL Protocol



GRPC



us-east-1b

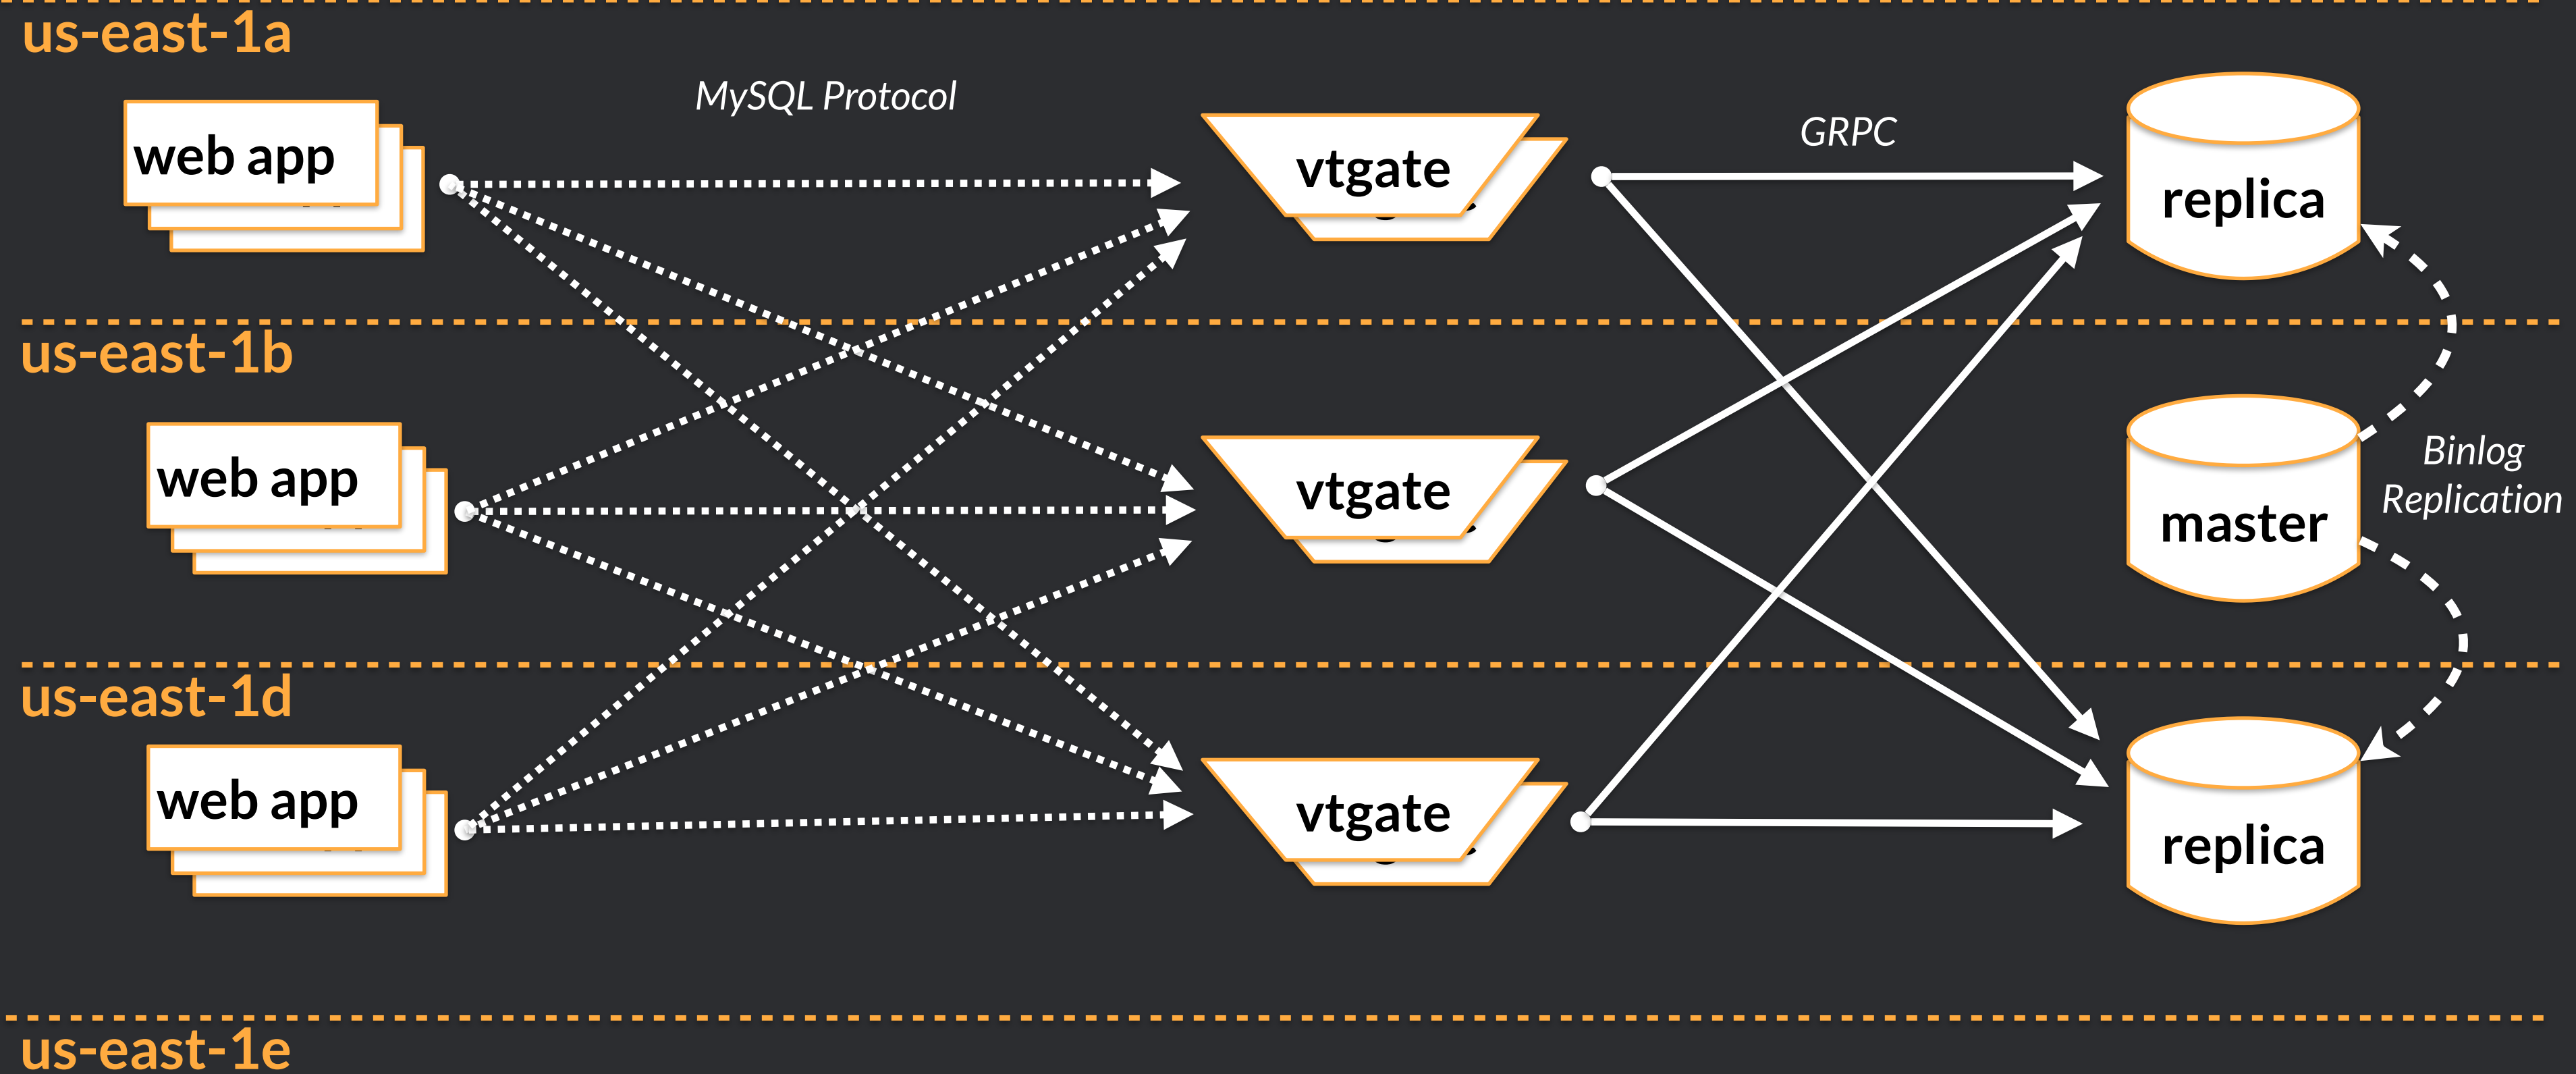


Binlog
Replication

us-east-1d



us-east-1e



AZ Aware Routing

us-east-1a



us-east-1b



us-east-1d



us-east-1e

Binlog
Replication


Improved... but still not great

Short-lived connections
require rapid open / close

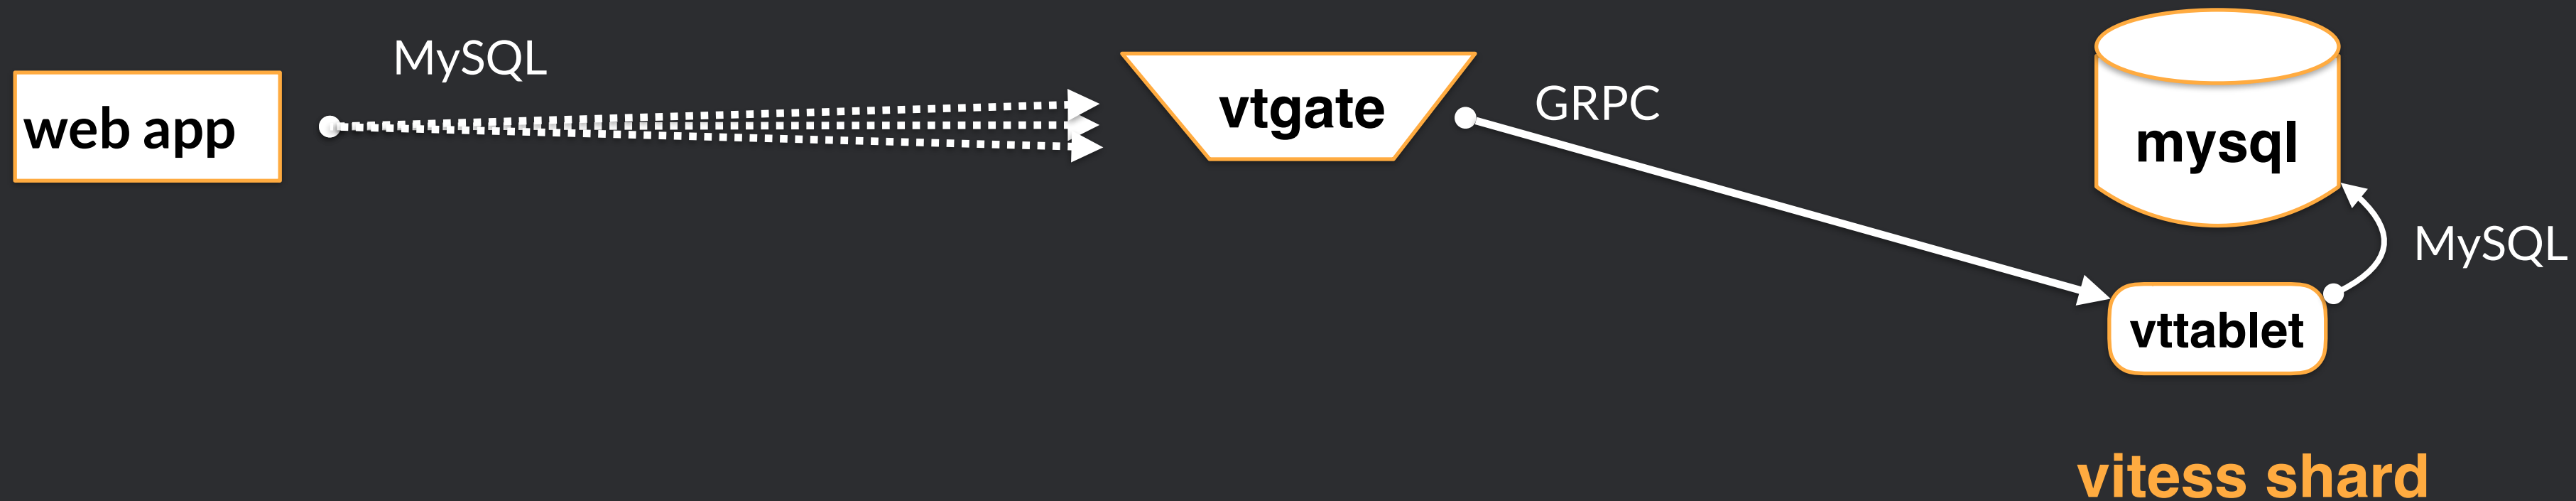
To mitigate packet loss, app
quickly fails over to try
another vtgate / shard

Under load this causes
delays, brownouts

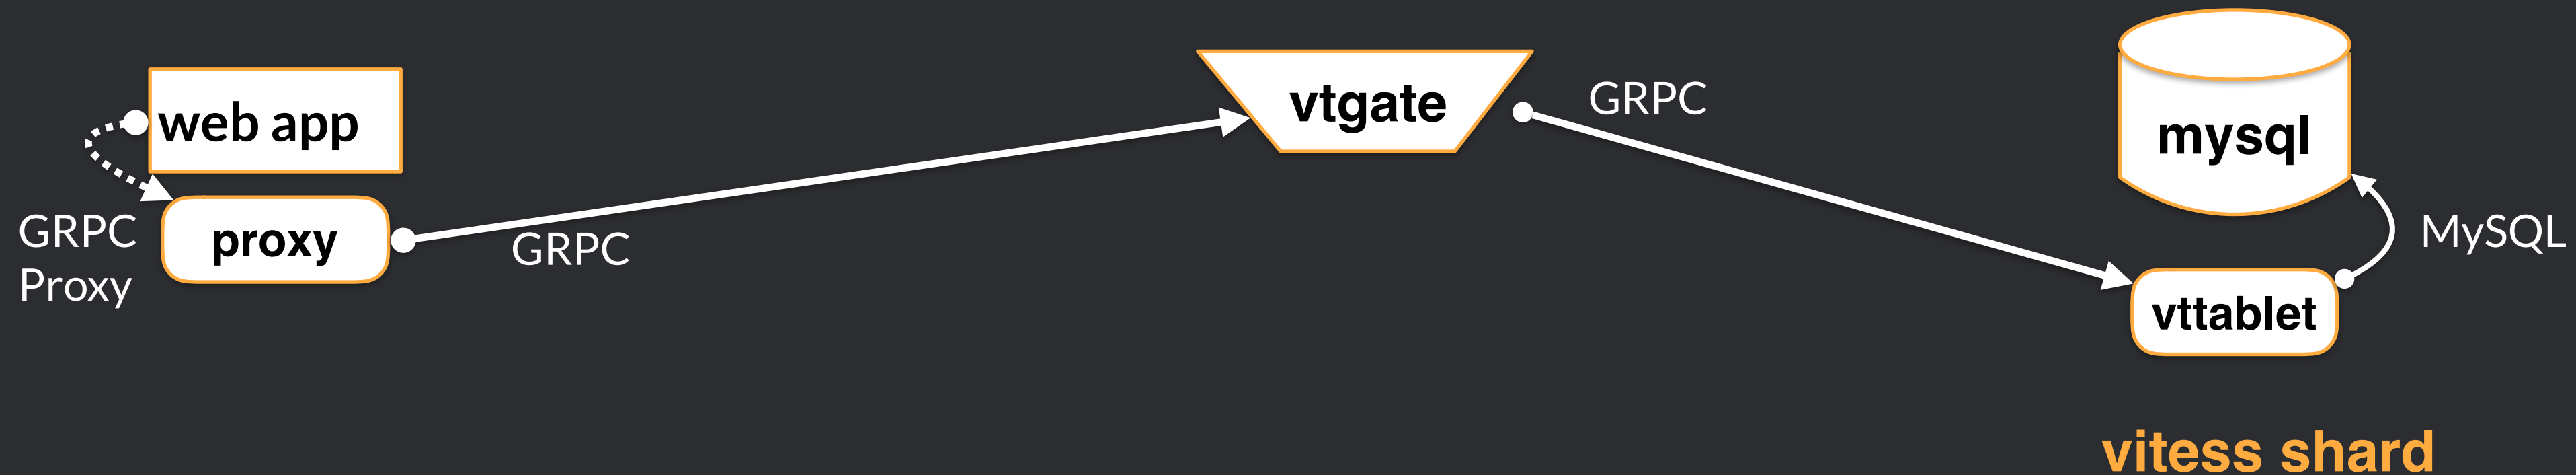
Long term goal: sticky
connections everywhere



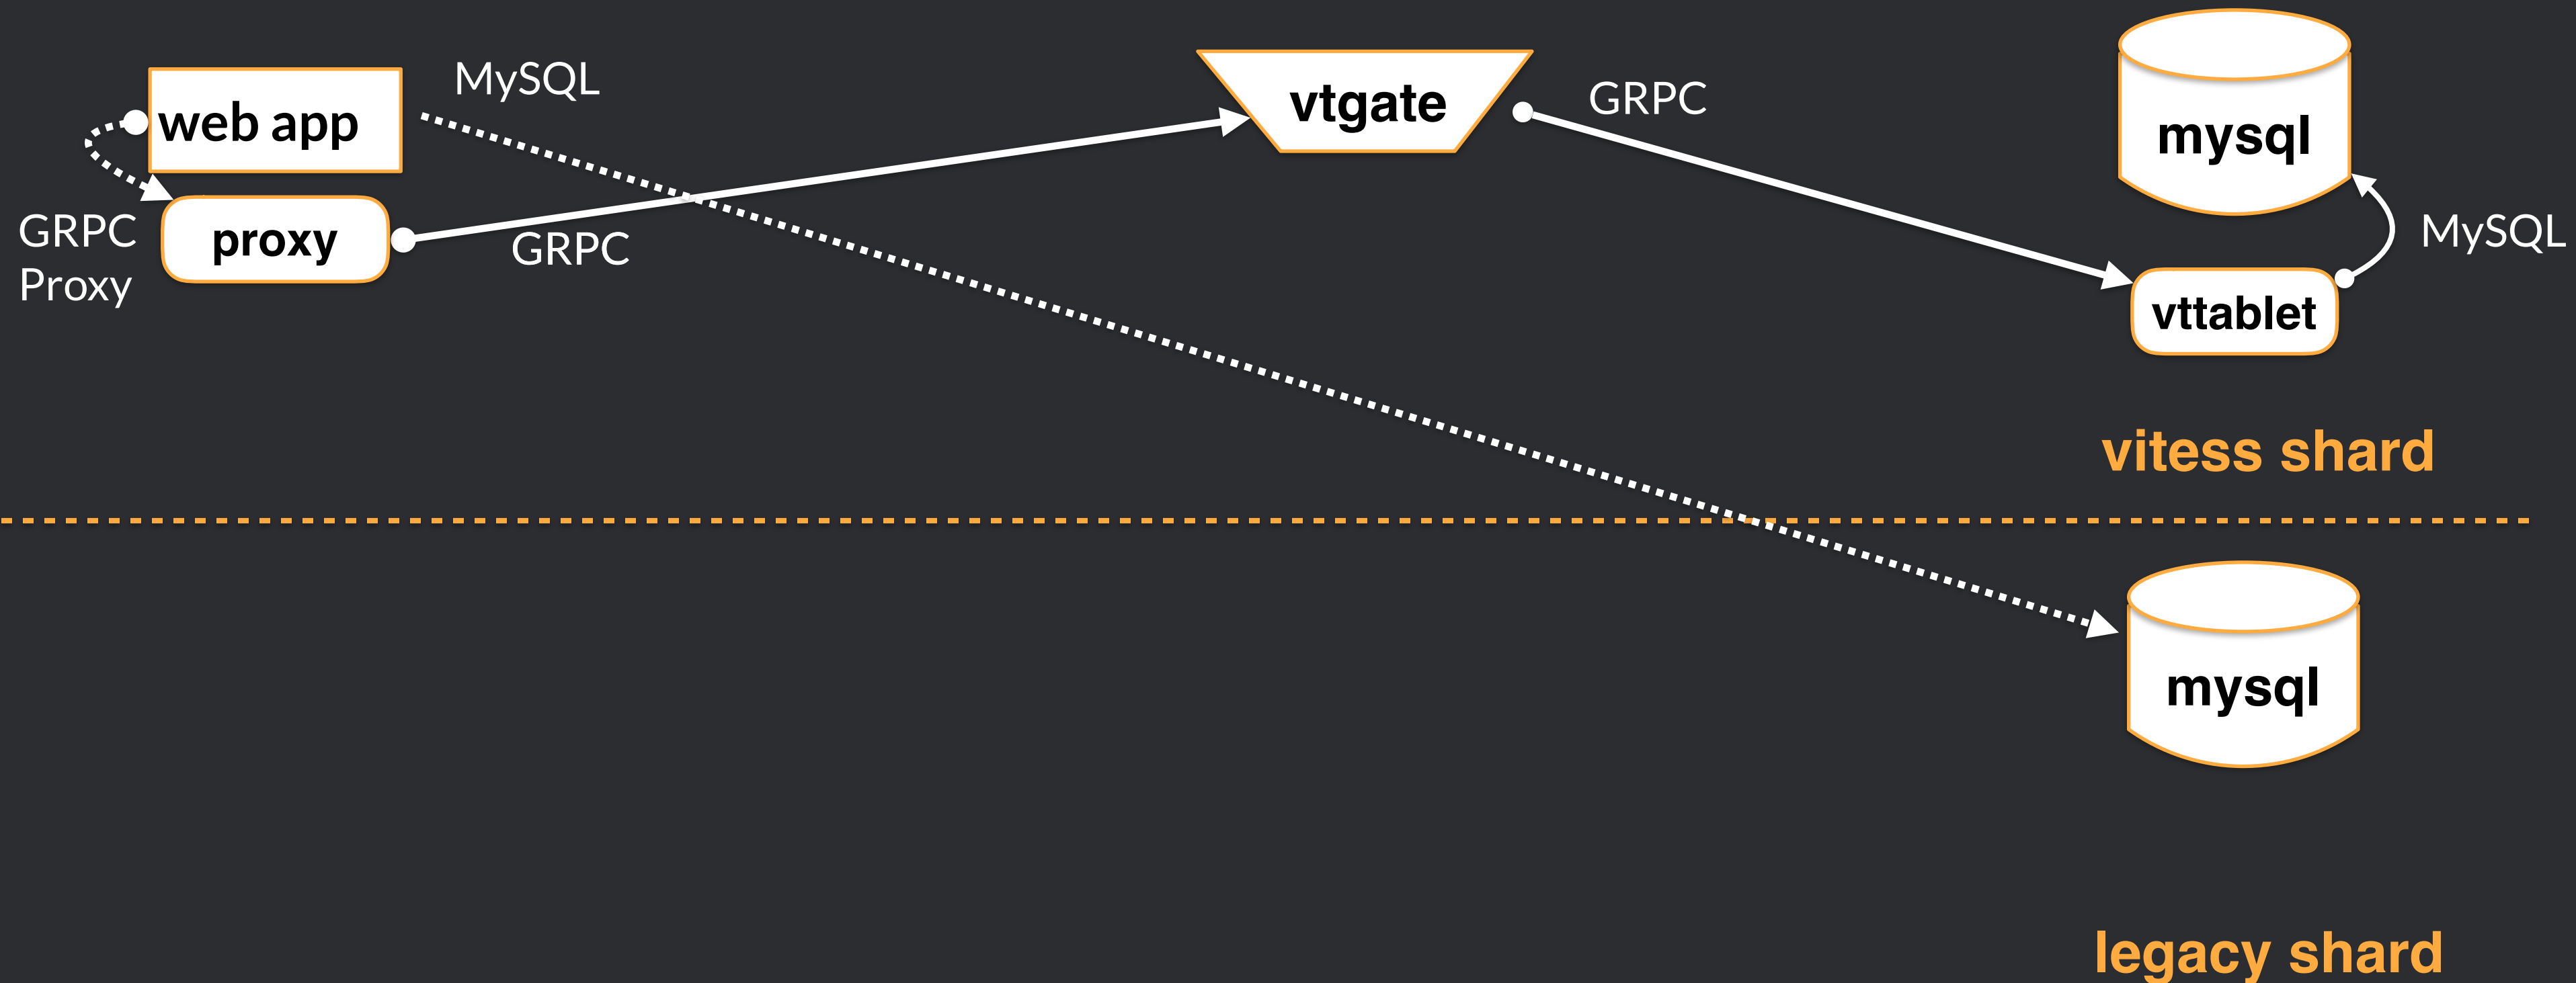
MySQL Connections



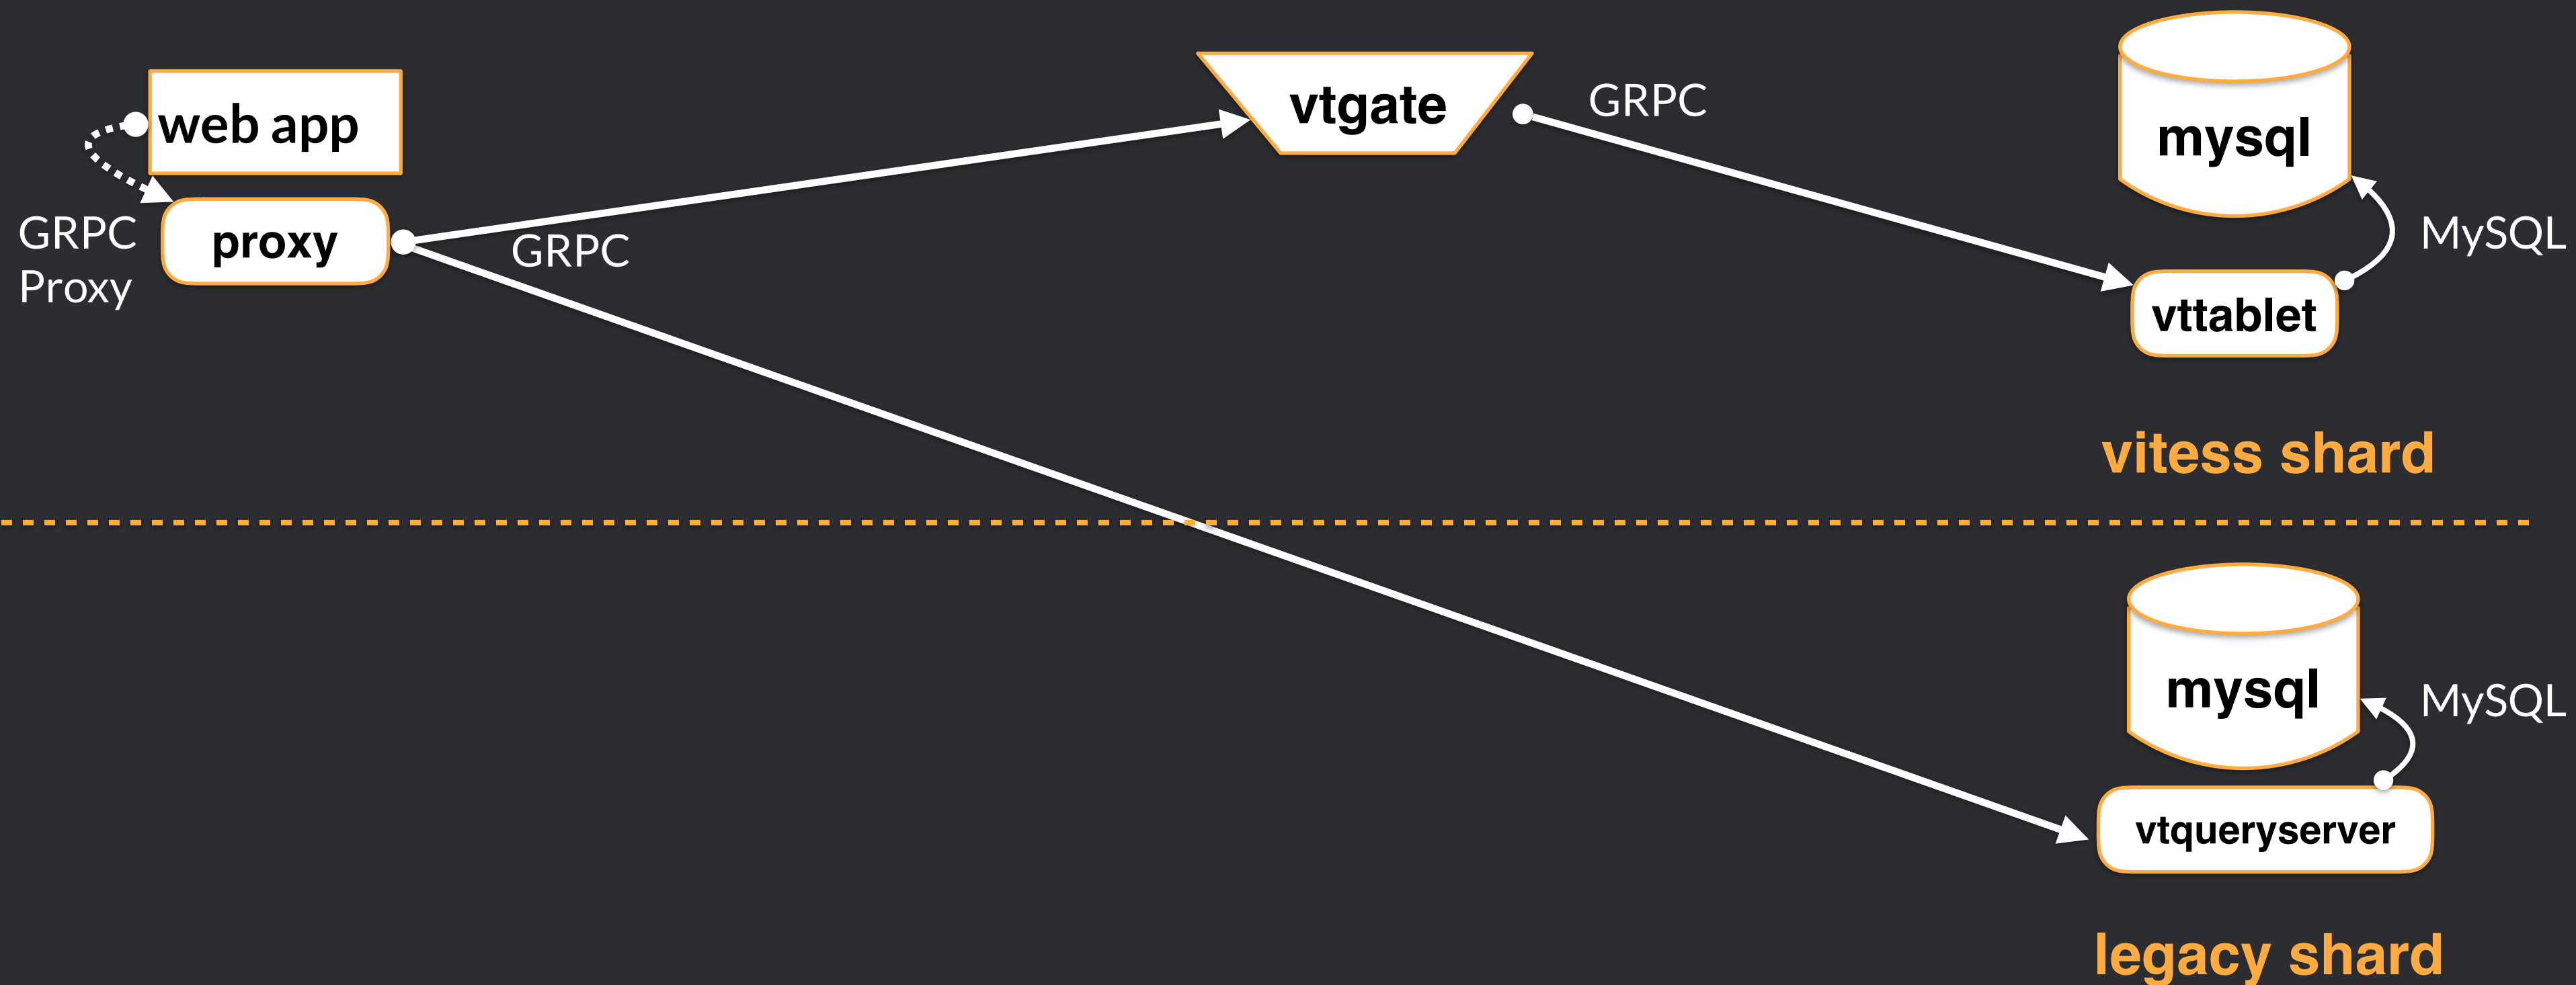
GRPC End to End



"Legacy" Databases



"Legacy" Databases (Future)



VTQueryserver Experiment

- Combine the vtgate query API (grpc + mysql) with the vtablet execution engine
- Helps protect mysql from query storms using connection pooling, hot row protection, query limits, etc
- Enables long lived GRPC connections from the web app
- Challenge to get the connection pool settings correct and to implement end-to-end prioritization

High Level Takeaways



Change All The Things

Because of Vitess, we had to:

switch to master / replica...

use semi-sync with gtid...

and orchestrator for failover...

But at the same time, we:

switched to row based replication...

on mysql 5.7 on new i3 EC2 hosts...

and an updated Ubuntu release...

using hhvm's async mysql driver...

and start reads from replicas...

Cha

📅 Wednesday 12:50PM-1:40PM

@ Room E

Designing and launching the next-generation database system @ Slack: from whiteboard to production

Guido laquinti - **Slack**

Slack is a messaging platform for teams that brings all communication together, creating a single unified archive accessible through powerful search.

MySQL is the primary storage for all our customer data and we currently execute billions of transactions per hours. As more users join the service, and Slack becomes a more critical part of their workflow, the system become more complicated and difficult to manage. What started out as a simple MySQL database was only the starting point for a long journey redesigning our entire database infrastructure.

This talk will analyze how our operations team took Vitess, a bleeding edge, poorly-documented open source software developed by Google and then hardened, tested and shaped it for our infrastructure and host all our mission critical data. This presentation will to thought the technical challenges that faced to successfully deploy this project (AWS instance upgrade i2 -> i3, storage SSD -> NVMe, kernel 3.13 -> 4.4, MySQL 5.6 -> 5.7, replication type async -> semi-sync, etc..) the key decisions that we took, what went well, what didn't and the course correction that we made along the way.

Attendees can expect to hear details about how we took some whiteboard conversations and turned them into battle-tested, production-caliber systems.

Beau

switch

use ser

and or

ation...

hosts...

ase...

iver...

...

Networking Matters

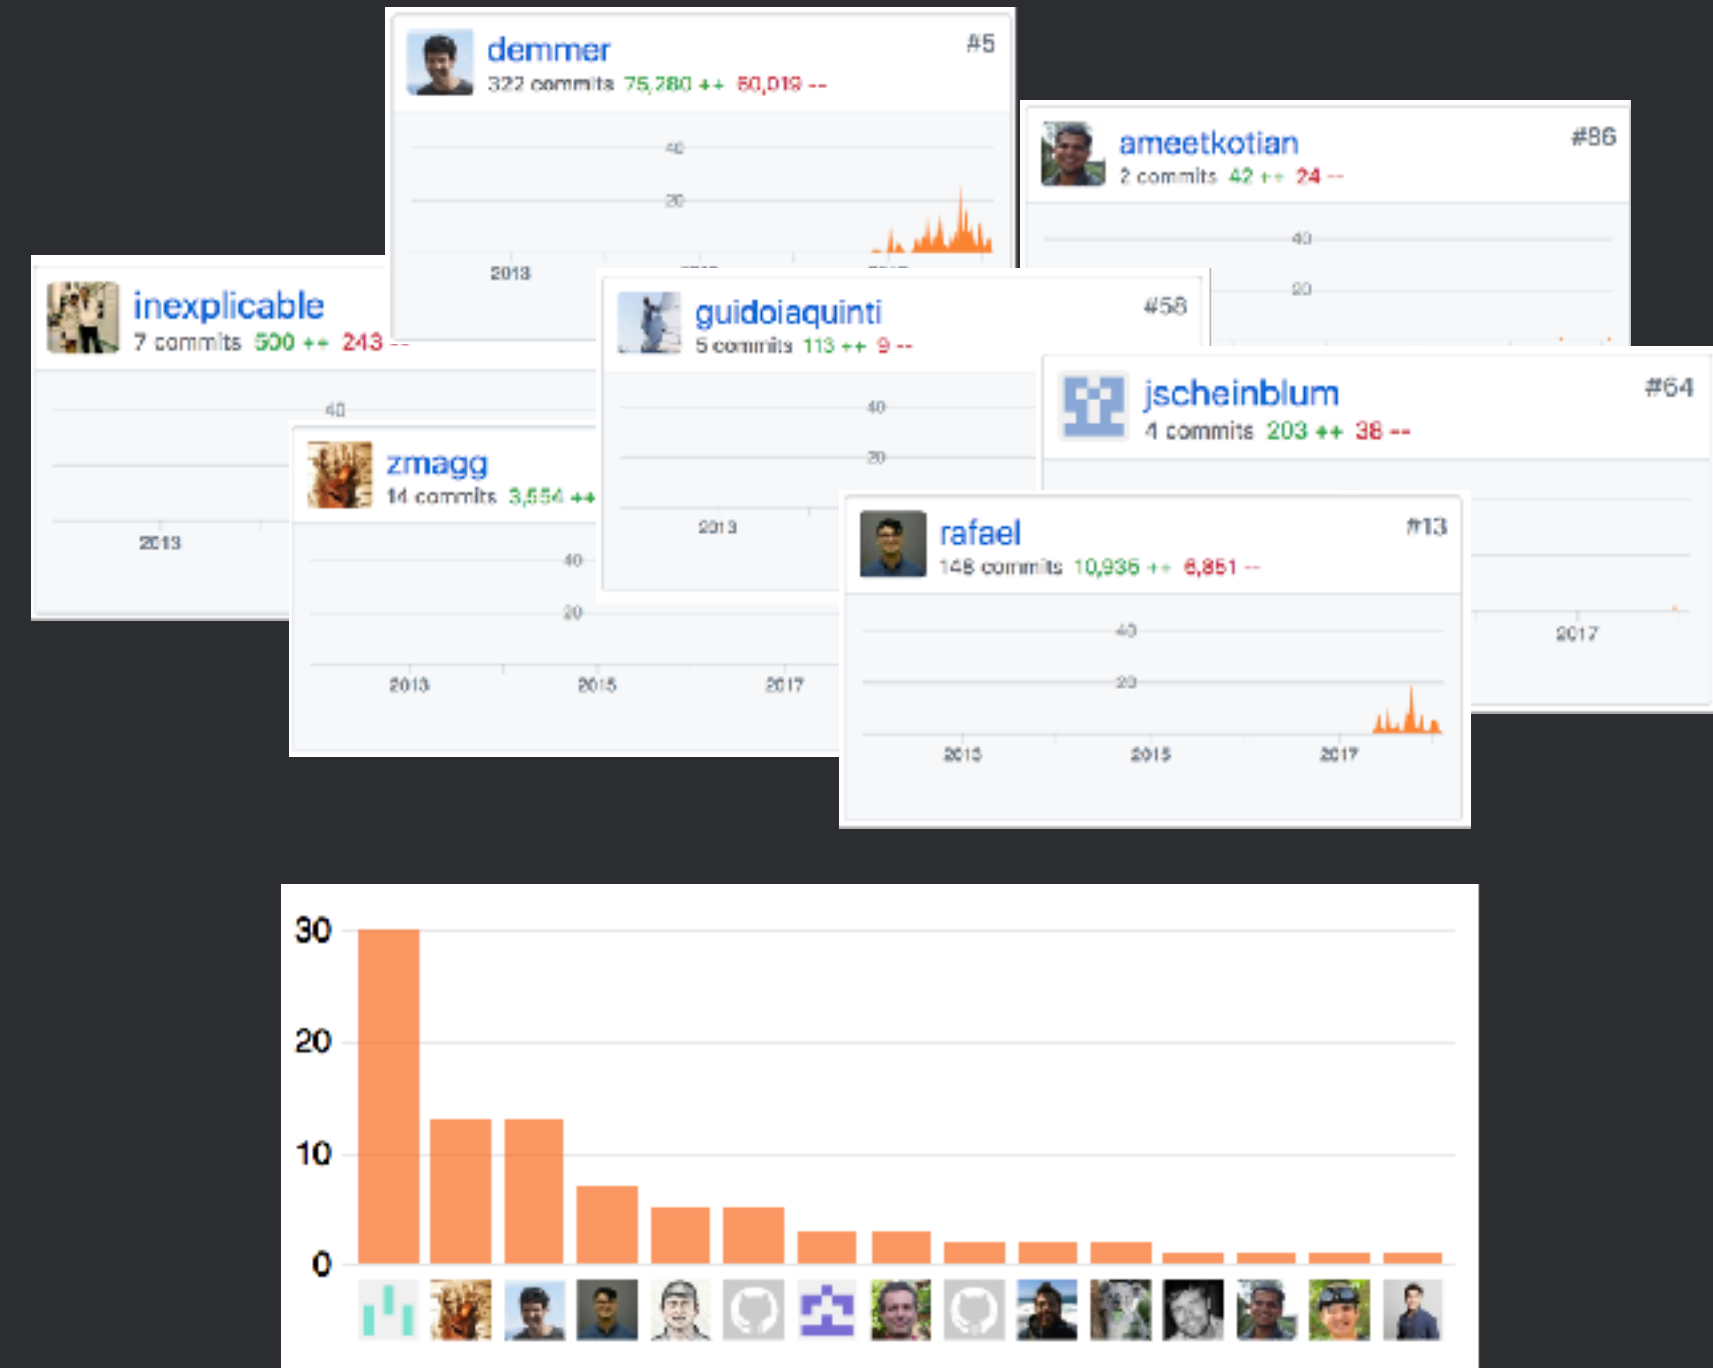
- Vitess is intrinsically more network dependent than our existing database architecture
- Performance depends (a lot) on network quality
- Improved consistency (single master / semi-sync) comes at the expense of availability and performance
- Able to work around some issues by kernel tuning, host placement, application routing to vtgate

Vitess: “Build” *and* “Buy”

The core of Vitess is stable, performant, and robust

But Slack’s use case differs from YouTube’s (and others)

Adoption required significant changes, all contributed back upstream



"Vitess is magical but not magic"

- !/? Besides MySQL, there are a still lot of new moving parts
- 😬 No ability (yet) to change sharding key
- 🚫 Still some unsupported queries (though not as many)
- ⚠️ Scalability / efficiency requires stale reads from replica
- 😞 Can't (yet) use familiar tools like phpmyadmin
- 🔍 Documentation!! -- many, many options to understand

Vitess At Slack: Thriving

- In production for ~10 months after ~7 months of effort
- Leadership buy in as the future for Slack databases
- Stable and performs well (so far)

We have a long but exciting road ahead...

 And we are hiring! 

Thank you!

